

THE OPERATOR ECONOMY

SMALL CORE. THICK PERIMETER.

A working thesis of small-core firms, rented capability, and
dependency-adjusted independence

The verdict

The original “Too Small to Bother” paper is a sharp manifesto wrapped around an unfinished theory. Its design looks resolved; its reasoning is not.

The phrase **Operator Economy is not completely unique**. It was used publicly in 2024 for a Web3 version of the gig economy [N1], in 2025 for AI-enabled fractional operators [N2], and now appears in unrelated leadership, commerce, and AI frameworks [N3–N7]. A separate Operator Class already offers an index, diagnostic, manifesto, archetypes, and an “Operator Economy” section [N4]. The exact phrase is therefore a crowded semantic field, not an uncontested invention.


The underlying ideas are also not new in isolation. Firm-boundary theory, outsourcing, vertical software, minimum efficient scale, company-of-one thinking, hidden champions, small giants, passion and creator economies, switching costs, and the claim that software can replace parts of a firm all predate this paper.

What can still be distinctive is the **specific causal model and measurement discipline** developed here:

- 1 permanent headcount is separated from total productive capacity;
- 2 capability moves from payroll into a rented perimeter;
- 3 lower production and coordination costs move scarcity toward demand, trust, judgment, compliance, and accountability;
- 4 apparent independence is discounted for channel, vendor, platform, and key-person dependence;
- 5 economic sufficiency is separated from competitive defensibility;
- 6 the theory publishes disproof conditions instead of calling preferences “axioms.”

That is a credible research program. It is not yet a proven new economy.

The strategic error to avoid is obvious: **a trademark cannot manufacture conceptual originality or market demand**. A trademark identifies the source of goods or services. It does not turn a familiar thesis into novel economic theory.



A FILING DOES NOT TURN THE LEFT-HAND BOX INTO THE RIGHT-HAND BOX

The personal truth

The repository shows a bias toward finished-looking artifacts: rubrics, pipelines, working papers, proof cards, scripts, carousels, and publication machinery. That strength has become a trap. It makes it possible to produce the appearance of intellectual closure before the category has a stable definition, denominator, dataset, or falsifier.

The only confirmed public embryo of this thesis matured at **394 impressions, two reactions, no comments, and no reposts**. That single result does not disprove the thesis. It does prove that the phrase and framing have not earned category demand. The later document post has no retained live URL, exact final copy, clean attachment, or performance record in the repositories.

Seeking legal ownership now would be a form of avoidance. It substitutes a filing event for the harder work: interview operators, build a coherent dataset, expose the model to counterexamples, publish corrections, and earn repeated use.

1. Source-state audit

There is no single canonical “Too Small to Bother” artifact

Four different artifacts have been blurred together:

Artifact	What the repository proves	What it does not prove
Public thesis embryo, week of 3 August	The florist/barbershop feed post published; tracker records 394 mature impressions and weak engagement	It did not carry the full working paper
Historical 10 August carrier copy	Recoverable only from brown-man-content commit 4552957; opens with a universalized business-survival claim	Exact final scheduler copy, live URL, and performance are absent
Retained reference PDF	Exists at studio/originate/too-small-to-bother/reference/working-paper-01-print-capture.pdf; content can be audited	It is not the clean public attachment
2,250-word manifesto/article	Exists in brown-man-content and contains the longest prose argument	Tracker says it did not publish; research.md incorrectly calls a different newsletter the published prose argument

The retained PDF is a broken Safari/iOS print capture:

- 38 physical pages representing a 19-page design;
- 18 blank or overflow pages;
- Claude Design URL, print timestamp, and browser furniture throughout;
- signature and source material spilling to physical page 38;
- metadata title “Manifesto PDF Design Brief”;
- explicitly labeled “REFERENCE ONLY, never ship” in the episode research pointer.

The tracker says a different clean 19-page export was scheduled. That clean file is absent from the relevant repositories and their referenced history. The historical carrier copy promises that every figure is sourced on the page where it appears; the retained PDF does not meet that standard.

This is not clerical trivia. A body of thought intended to compound needs exact public artifacts, dates, URLs, versions, hashes, corrections, and performance records. Without them, later claims of priority, authorship, iteration, or public use become difficult to prove.

Immediate provenance repair

- 1 Recover the live LinkedIn document-post URL and lifetime metrics from LinkedIn.
- 2 Download the exact public PDF and final caption from the live post if still available.
- 3 Store the bytes unchanged, record SHA-256, publication time, URL, and screenshots.
- 4 Correct research.md so it does not label the unpublished manifesto as published.
- 5 Keep the broken print capture as a reference artifact, never as a public master.
- 6 Create a version ledger for every future working paper: draft, approved, exported, published, corrected, superseded.

2. The category claim

Working definition

An **operator firm** is an enterprise in which a small, owner-controlled permanent core retains:

- the customer promise;
- the right to set the offer, price, and allocation of resources;
- the acceptance standard for the work;
- customer-facing commercial accountability and recourse;

while assembling a material share of its execution capacity from external capabilities:


- software and AI systems;
- contractors and fractional specialists;
- platforms and marketplaces;
- logistics, payments, capital, and cloud infrastructure;
- licensed, regulated, or otherwise institutional partners.

The **Operator Economy** is the working name for a hypothesis: old networked-firm structures are diffusing into more sectors and becoming feasible at a smaller permanent-core size because external capability can be bought in finer units. The form itself is not new. The empirical claims are wider diffusion, a lower feasible core in compatible work, and a changed distribution of control, risk, and value.

This is a boundary configuration, not a headcount slogan. “One person” is neither necessary nor sufficient.

Plain-English version: A small permanent core owns the customer promise and recourse. A larger perimeter supplies capacity and retains explicitly mapped duties and liabilities.

CAPACITY SITS ACROSS THE PERIMETER. CONSEQUENCE STAYS AT THE CORE.



Operational boundary for the first dataset

Words such as “small,” “material,” and “owner-controlled” are loopholes unless they have working rules. Use these rules for the first empirical cohort:

Term	Working rule
Countable unit	One firm × one materially distinct offer or business line × one calendar year; classify the same entity again when its role or boundary changes materially
Permanent core	Working owners plus any person, regardless of payroll or contractor label, who performs at least 0.5 full-time-equivalent firm-specific work for six continuous months and cannot be substituted through the firm's normal supplier process
Small core	No more than ten full-time-equivalent people under the permanent-core rule; this is a research cutoff to test, not a natural law
Owner-controlled	Working owners collectively hold more than 50% of ordinary voting or equivalent residual control rights and, within legal and contractual constraints, can accept or reject counterparties and retain final authority over offer, price or payer terms, cash allocation, acceptance, and customer recourse
Material capability perimeter	At least one external capability either accounts for 20% or more of recurring economically necessary delivery cost, excluding raw materials, rent, and commodity utilities; or its removal would interrupt more than 20% of delivery or revenue for longer than 30 days without a tested substitute
Accountable outcome	A named legal entity and human decision-maker give the buyer recourse and do not disclaim the core customer promise to a platform, model, contractor, or marketplace; professional and regulatory duties elsewhere in the perimeter remain separately attributed

These thresholds prevent payroll relabeling. A nominal contractor who works like a permanent team member counts inside the core. A business does not qualify merely because it uses email, cloud hosting, a bookkeeper, or ChatGPT. The external perimeter must carry an economically material or operationally critical capability.

The 10-person, 20%, six-month, and 30-day thresholds are preregistered starting points. Field evidence may move them, but any revision must be published before outcome analysis, not chosen afterward to rescue the thesis.

Classify roles by transaction, not by flattering identity. A capability supplier can itself be an operator firm in a different firm-offer-year unit. Record each purchase once in the buyer's perimeter ledger; do not add shared-vendor capacity to multiple buyers as if each owned it.

Why “economy” is still provisional

An economy label earns its keep when it identifies:

- 1 a stable unit;
- 2 a recurring transaction or resource flow;
- 3 a causal mechanism;
- 4 a measurable population and denominator;
- 5 predictions about who captures value and who bears risk;
- 6 boundaries that prevent the term from swallowing everything.

This draft supplies a working unit and hypotheses. It does not identify a transaction unique to operator firms: customers buy an outcome and recourse from the core, while the core buys modular capabilities from the perimeter—an exchange pattern long used by general contractors, film productions, publishers, brokerages, and other networked organizations.

The claim therefore stands or falls on changed prevalence and minimum feasible core size, not invention of the form. Until original data shows the classifier is reliable, the configuration is rising across repeated periods, and its control and risk profile

differs from matched conventional small firms, “Operator Economy” should be described as a **media identity, working category, and research program**, not a statistically established sector.

Promotion to an empirical economy claim requires all of the following:

- 1 independent coders classify at least 100 cases with agreement of 0.80 or better;
- 2 a probability-based or defensibly weighted sample estimates prevalence with uncertainty in at least three defined sectors and two periods;
- 3 the lower confidence bound exceeds 5% of eligible firm-offer-year units in those sectors and prevalence rises materially across the periods;
- 4 matched comparisons reveal at least one preregistered difference in boundary, risk incidence, or performance that ordinary small-firm status alone does not explain;
- 5 another research group can reproduce the coding and headline results.

Those thresholds are governance rules for this research program, not universal laws of what counts as an economy.

The unit is the boundary, not the person

The original paper treats payroll headcount as productive capacity. That is its most consequential mistake.

A two-payroll-employee telehealth company can sit on physicians, pharmacies, prescribing infrastructure, payment rails, cloud systems, counsel, compliance, capital, and distribution. Calling it a “two-person company” describes a legal payroll boundary. It does not describe the productive system.

The useful question is not “How many employees?” It is:

Which capabilities are held inside the permanent core, which are rented, who controls each one, and who remains accountable when it fails?

That boundary map is the proper unit of analysis.

3. What qualifies — and what does not

Adjacent category	Primary unit	Who normally controls the customer promise?	Relationship to an operator firm
Gig worker	Task or shift mediated by a platform	Platform or buyer	Usually excluded when the platform sets price, rules, access, and recourse
Freelancer	Individual skilled labor	Often the client, sometimes the freelancer	Can become an operator firm when the person owns the offer, customer relationship, repeatable system, and P&L
Fractional executive	Skilled capacity sold across clients	Client	Adjacent; not automatically an operator firm because the practitioner may still sell labor inside clients' systems
Creator	Audience and content	Creator or platform	May qualify when the audience is a distribution asset inside a broader operator-controlled firm
Agency	Coordinated client service	Agency	Can qualify; headcount and service work do not disqualify it
Traditional small business	Local or vertical enterprise	Owner/firm	May qualify if its boundary and control structure fit; technology is not required
Venture startup	Growth-oriented company	Company, board, and investors	Qualifies only if working owners retain the residual-control rule; funding alone does not decide inclusion
Agentic economy	Programmatic transactions among agents	Unsettled across users, firms, and platforms	Distinct: the Operator Economy keeps a human or legally accountable firm at the center
Platform business	Market or infrastructure intermediary	Platform	Often a capability supplier or landlord to operator firms, not the operator firm itself

Inclusion test

An enterprise belongs in the working category when all four conditions are present:

- 1 **Promise ownership:** the core makes and owns the customer promise.
- 2 **Economic control:** the core controls the offer, price, cash, and resource allocation.
- 3 **Capability assembly:** the external perimeter passes the materiality or criticality rule defined above.
- 4 **Outcome accountability:** a named human or firm remains answerable for the result; AI and vendors do not become accountability theater.

This test deliberately excludes “used ChatGPT once,” a contractor completing someone else's task, and a platform worker mislabeled as an entrepreneur.

4. The boundary-shift mechanism

The strongest version of the thesis is not “AI made building cheap.” That statement is too narrow, historically late, and frequently false.


The deeper change is that many forms of institutional capability can be purchased in smaller increments and coordinated through software:

- compute by request rather than a data center;
- payment and identity checks by transaction;
- counsel, finance, design, and operations by project or fractional allocation;
- logistics and fulfillment by shipment;
- models by token or API call;
- specialized regulated capacity through licensed partners;
- demand through marketplaces, affiliates, or paid channels.

Cloud services and vertical software were already shifting firm boundaries before generative AI. AI adds divisible cognitive capacity, but it is one layer in a longer modularization of the firm.

The causal chain

- 1 **Capability can become more granular.** The smallest purchasable unit gets smaller in tasks with usable interfaces, standards, and verification.
- 2 **Some fixed commitments can become variable commitments.** Payroll, infrastructure, and capital may become metered operating expense, but integration and minimum contracts can preserve fixed cost.
- 3 **The permanent boundary can contract under moderators.** Asset specificity, learning loss, security, regulation, bargaining exposure, and integration cost can instead force capability back inside.
- 4 **Coordination and verification must land somewhere.** The operator, a platform, or a regulated partner specifies, sequences, checks, and accepts outputs; the paper must measure who actually does this work.
- 5 **The bottleneck moves.** Demand, trust, judgment, compliance, context, and accountability become relatively scarcer.
- 6 **Risk is reconfigured, not erased or necessarily exchanged one-for-one.** Payroll and capital exposure may fall while vendor, platform, data, labor-incidence, regulatory, and key-person exposure can be added.
- 7 **Value capture is contested.** Scarcity matters, but market power, property rights, contracts, regulation, and switching costs determine whether the operator, worker, platform, incumbent, or capability supplier captures the gain.



The key correction

The original paper says both:

- build cost broke the old Team → Capital → Outcome → Market chain; and
- build cost was never binding because distribution was.

The reconciliation is:

Production and coordination costs determine which offers are feasible. Demand, trust, and accountability determine which feasible offers survive.

Both can bind at different stages. There is no universal single bottleneck.

5. Seven propositions — with disproof conditions

These are propositions, not axioms. An axiom is assumed. A serious thesis invites a test.

Proposition 1 — Capability compression is conditional

In sectors where tasks are digitally representable, interfaces are standardized, and verification is affordable, software, AI, and external specialists can reduce the permanent capacity needed to operate.

It does not imply: every job can be automated, every firm should shrink, or one person can perform the total work of a larger institution.

Disproof or boundary evidence: no decline in fixed commitments or permanent core size after controlling for sector, firm age, outsourcing, and demand; integration and verification costs erase the apparent saving.

Proposition 2 — Headcount increasingly understates productive capacity

As more capability sits with vendors, platforms, contractors, and infrastructure, payroll employment becomes a weaker proxy for the resources used to produce an outcome.

Disproof or boundary evidence: external capability share remains immaterial, or payroll-adjusted and full-system capacity measures produce the same conclusions.

Proposition 3 — Scarcity migrates; it does not disappear

When production becomes cheaper and more widely available, scarce buyer attention, trust, workflow access, distribution, judgment, and accountability become more important.

Disproof or boundary evidence: broad access to generic production tools produces durable, above-normal returns without differentiated demand access, context, trust, or integration.

Proposition 4 — Generic tool access is symmetric

If the same model, template, or no-code stack is available to entrants, peers, platforms, and incumbents, access to it is not a moat. Advantage must come from what compounds around the tool: direct demand, permissioned context, workflow position, proprietary learning, judgment, brand, or controlled assets.

Disproof or boundary evidence: generic tool access alone predicts persistent performance after controlling for distribution, experience, integration breadth, and sector.

Proposition 5 — Rented scale creates a dependency paradox

The operator can gain capacity without adding payroll while becoming more dependent on revocable external systems. Apparent independence can therefore increase at the ownership layer and fall at the operating layer.

Disproof or boundary evidence: outsourced capacity is readily substitutable, portable, and low-concentration, with no meaningful outage, repricing, data, or policy exposure.

Proposition 6 — Control concentration creates speed and fragility

A concentrated operator core can decide and adapt quickly. The same concentration can create key-person risk, quality bottlenecks, weak transferability, burnout, and low resale value.

Disproof or boundary evidence: concentrated decision rights do not predict service interruptions, owner overwork, customer risk, or transfer discounts; documentation and delegated controls eliminate the effect.

Proposition 7 — Sufficiency is real; “too small to bother” is not a moat

A narrow, reachable segment of a large and costly problem can be sufficient for one accountable operator while remaining below the standalone entry hurdle of some scaled competitors.

That condition can deter a large de novo entrant. It does not deter:

- another small operator;
- a platform bundling the feature;
- an adjacent incumbent;
- a customer building internally;
- a roll-up or acquirer;
- a clone with better distribution.

Disproof or boundary evidence: the firm cannot generate normalized economic surplus inside its chosen capacity; small peer entry, bundling, or platform action destroys the economics; “smallness” fails to predict competitive attention.

6. The evidence ledger

What current evidence supports

Evidence	What it supports	What it does not support	Confidence
Census SUSB employer-business methodology [E1]	The legacy software-publisher counts describe firms or establishments with paid employees	Solo firms, total humans involved, contractors, or the size of an Operator Economy	High for universe definition
Legacy paper's 2017–2022 SUSB calculation [R1]	A rise in small employer software publishers in that NAICS series	AI causation, durability, profitability, median outcomes, or all software companies	Moderate pending reproducible data package
Census Nonemployer Statistics [E2]	A separate universe exists for tax-filing businesses with no paid employees	A clean join to SUSB without classification, cutoff, and comparability work	High
2026 Census BTOS [E3]	Overall AI use was 17–20%; under-20-employee adoption did not change significantly in the observed six months; fewer than 20% of firms with four or fewer employees reported use	Small firms leading an economy-wide AI transformation	High, descriptive
2026 Census working paper [E4]	18% of firms used AI; 66% of adopters used it only to augment; AI-linked employment decreases appeared in 2%; integration breadth correlated with commercial performance	Causal proof that AI creates one-person firms or eliminates teams	High for reported survey; correlation is not causation
OECD competition review [E5]	AI can lower minimum efficient scale in some tasks and sectors	Even or automatic benefit to small firms; the review also reports much higher adoption by large firms and platform dependencies	Moderate to high synthesis
Marchesi and Tang working paper [E6]	AI-compatible industries experienced more than 10% higher firm formation after ChatGPT in its difference-in-differences design	Settled causal fact across all sectors; the paper is an emerging working paper, not final consensus	Provisional
NBER firm-size and boundaries work [E7]	Off-the-shelf cloud services can reduce fixed cost and minimum efficient scale	A uniquely generative-AI mechanism	Established lineage
App and subscription benchmarks in Working Paper No. 01	Outcomes are highly skewed in selected app/subscription samples	A time-series decline in the median operator firm, or one common population across apps, firms, creators, and establishments	Illustrative only

Evidence that directly challenges the hype


The current data does not support “small firms are winning because AI.”

- Census reported less than 20% AI use among firms with four or fewer employees in the period ending May 2026 [E3].
- AI use rose in firms with at least 20 employees while the change among smaller firms was not statistically significant during the six-month observation window [E3].
- Census researchers found most adopting firms used AI for augmentation, not headcount removal [E4].
- OECD reported a 3.3-fold large-versus-small adoption gap across its 2024 area data and warned that data, compute, integration, and lock-in can reinforce incumbents [E5].

The defensible conclusion is conditional:

AI can reduce the minimum feasible core in compatible workflows, but the ability to integrate it, verify it, access data, carry liability, and reach buyers determines who captures the gain.

That conclusion is less viral and more likely to survive contact with evidence.



7. Claims from Working Paper No. 01 that must be killed or repaired

Original claim or move	Verdict	Repair
"Most software companies in America now have four people or fewer"	Kill	"In the cited 2022 SUSB series, 57.5% of employer software-publisher firms were in the 1–4 paid-employee class." Owners, contractors, and nonemployers are not counted [R1; E1–E2].
2017–2022 count growth proves the post-AI floor collapse	Kill	Treat it as evidence of a longer cloud, modularity, and outsourcing trend that AI may accelerate. Test the post-2022 effect separately.
The firms are "durable"	Kill until measured	Use cohort survival, churn, ownership duration, and normalized economic surplus. Two stock counts are not durability.
\$674,676 describes the typical tiny firm	Kill	It is a derived arithmetic mean of aggregate receipts, not a median, profit, or owner income.
A wider top-to-bottom ratio means the median got worse	Kill	A ratio can widen because the top rose. A time series of medians or distribution quantiles is required.
One person now does what three to five did	Kill or make task-specific	Measure particular workflows, output quality, time, and external labor. Do not generalize from anecdotes.
Medvi is a two-person productive system	Kill	Use it as a boundary case: a tiny payroll core orchestrating a thick regulated perimeter.
Team → Capital → Outcome → Market is a chain that broke	Reframe	Each link is conditional. Capital, obligations, and external institutions remain; some fixed costs became variable.
Build cost was never binding; distribution was	Reframe	Production governs feasibility; demand and trust govern survival; compliance and accountability can bind both.
Fluency is the entry ticket, not the moat	Keep with caveat	Fluency helps specification and verification. Outsider firms can acquire it; retained learning and position matter more.
"Your protection and your ceiling are the same number"	Keep only as provocation	Separate the operator's chosen capacity ceiling from competitor attention, peer entry, platform bundling, and acquisition.
"Too small to bother" is defensibility	Demote	It is one possible deterrent against one entrant class, not a general moat.
"Two thirds of everything ever started" disappears in a decade	Kill	The cited BLS evidence describes one historical establishment cohort, not everything ever started.
Every figure is sourced on the page where it appears	False for retained PDF	Use full on-page locators or endnotes with titles, dates, URLs, and claim-level mapping.

8. The boundary-ledger method

The prescriptive method has four layers and five ledgers. The label is deliberately descriptive, not a claim to an ownable name. It should remain a structured assessment, not a weighted index, until real outcome data can justify weights.

Layer A — The operator core

Capabilities that should remain controlled by the accountable core:

- customer promise and recourse;
- offer and pricing;
- resource allocation and cash authority;
- acceptance standards and irreversible judgment;
- risk decisions and escalation;
- relationship memory that cannot safely be outsourced.

This does not mean the operator personally performs every task. It means the operator cannot outsource accountability while continuing to claim control.

Layer B — The capability perimeter

For every external capability, record:

- supplier and commercial model;
- data or system access;
- decision authority;
- verification method;
- replacement time;
- price-change exposure;
- outage or revocation impact;
- regulatory or contractual accountability;
- fallback.

The perimeter includes humans and institutions, not merely AI tools.

Layer C — The bottleneck

At any moment, identify the actual scarce constraint:

- feasibility and production;
- buyer access;
- trust and proof;
- operator judgment;
- delivery capacity;
- compliance or liability;
- cash timing;
- platform permission.

The bottleneck can move as the business grows. “AI makes building easy” is useless if demand, verification, or liability is binding.

Layer D — The control ledger

Track five dimensions without collapsing them into a fake precision score:

Ledger	Core measures	Question
Economics	Normalized economic surplus; minimum viable customer count; reachable buyer pool; cash conversion and reserves	Is this a business after paying for labor, capital, and risk?
Control	Offer and price authority; direct customer identity; data portability; cash and payment control	Which economic rights can a platform or client revoke?
Resilience	Largest dependency blast radius; replacement time; customer concentration; key-person exposure	What single failure can stop delivery or revenue?
Position	Retention; repeat purchase; workflow depth; learning accumulation; time and cost to reproduce; entrant-class threat map	What improves with each customer or operating cycle?
Incidence	Contractor pay and volatility; benefits and classification; unpaid family labor; training burden; supplier labor conditions; public or customer risk transfer	Who absorbs the obligation or downside removed from the permanent payroll?



LEDGER FIRST. WEIGHTS ONLY AFTER OUTCOME CALIBRATION.

Accountability is not one field. Map at least five separate layers:

- 1 customer recourse for the promised outcome;
- 2 contractual liability between the core and every supplier;
- 3 nondelegable professional duties held by licensed people or institutions;
- 4 regulatory liability assigned by law;
- 5 residual loss-bearing when contracts and insurance do not make the system whole.

The core does not “own every consequence.” It owns the customer-facing promise and must disclose, coordinate, and verify the accountability map. A named human must retain approval over irreversible commitments, final acceptance, risk escalation, and customer recourse; licensed professionals retain decisions the law does not allow the core to absorb.

Dependency-Adjusted Independence

Legal ownership is not the same as operating independence. A founder may own every share while a platform controls demand, a payment processor controls cash access, a model vendor controls core production, or one licensed partner controls delivery.

Dependency-Adjusted Independence is the working term for independence discounted by critical outside revocability, concentration, and replacement friction.

At this stage it is a structured assessment, not a measured discount, index, or score. Do not claim quantitative independence until independent coders can reproduce the assessment and its components predict observed interruption, recovery, or survival.

Do not turn it into a score yet. For every dependency that can stop revenue or delivery, record:

- the share of demand, capacity, cash, data, or legal authority exposed;
- who holds the right to revoke, reprice, suspend, or restrict it;
- time to detect failure and time to replace the dependency;
- data and workflow portability;
- whether a fallback has been tested under realistic conditions;
- the maximum interruption the firm can survive.

A firm is not operationally independent merely because alternatives exist on a vendor list. It is independent only to the extent that critical rights and capabilities can be retained or restored inside its survival window.

Normalized economic surplus

Cash profit overstates a founder-dependent firm's economics when the owner works for free or carries unpriced risk.

Use mutually exclusive cost buckets:

- **third-party operating cash costs** include paid employee and contractor labor but exclude owner compensation, owner distributions, capital contributions, and transfers into reserves;
- **owner and unpaid-family labor charge** uses documented hours and a preregistered replacement-cost benchmark, replacing rather than supplementing any owner compensation line;
- **owner-supplied capital charge** records the opportunity cost of capital actually at risk, using a stated benchmark and sensitivity range;
- **current-period reserve contribution** covers deferred maintenance, business taxes not already in operating costs, and risk capital not already funded.

Normalized economic surplus = revenue – third-party operating cash costs – owner and unpaid-family labor charge – owner-supplied capital charge – current-period reserve contribution.

This is not an accounting standard. Every bucket, benchmark, and sensitivity range must be published. It is an anti-self-deception measure, not a valuation claim.

A right-sized business that depends on permanent underpayment of its owner or unpaid family labor is not economically sufficient. It is a subsidized job.

Reachable market, not theatrical TAM

Define the reachable buyer pool as buyers the operator can credibly contact, earn access to, or acquire through tested channels inside a fixed time and payback budget.

The relevant question is not “How large is the market?” It is:

How many buyers can this operator actually reach, at what cost, with what proof, before cash or attention runs out?

9. The viability gate

The old paper conflates size, defense, ambition, and success. This gate separates them. It is a break-even sanity check, not a distinctive economic theory or proprietary name.

Let:

- **V** = minimum serviced demand required for normalized economic surplus to reach zero; the reserve contribution is already inside that calculation;
- **R** = demand realistically reachable through tested channels;
- **C** = maximum demand the chosen operating design can serve without breaking quality, control, or resilience.

For one defined offer, demand unit, price, product mix, churn assumption, and time period, an operator firm passes the viability gate when:

V is less than or equal to the smaller of R and C.

This says the firm can reach enough demand to be economically sufficient without exceeding its chosen design. V, R, and C are not comparable until those definitions and uncertainty ranges are fixed.

It says nothing about defensibility.



DEFENSIBILITY IS A SEPARATE THREAT SCREEN — SMALLNESS IS NOT A MOAT

Defensibility is a separate threat screen

Evaluate at least five entrant classes:

Entrant	Why “too small” may fail	Relevant response
Peer operator	Needs far less reward than a scaled company	Direct demand, faster learning, reputation, workflow position, service quality
Adjacent incumbent	Can cross-sell to existing customers	Integration depth, focused service, neutral interoperability, trusted niche position
Platform	Can bundle, reprioritize, tax, or remove the category	Customer and data portability, multi-channel fallback, contractual protection
Customer insourcing	Buyer can rebuild when value becomes visible	Continuous learning, lower total cost, hard-to-coordinate perimeter, shared risk
Roll-up or acquirer	Small firms can be attractive in aggregate	Decide whether acquisition is acceptable; do not confuse acquisition interest with competitive entry

The strongest retained idea from the original is a rate comparison:

A position strengthens when accumulated customer context, trusted access, workflow fit, and learning grow faster than an entrant's ability and incentive to reproduce them.

That is a hypothesis to measure, not a universal law. Embeddedness can also become customer-hostile lock-in. A thesis about ownership should not excuse captivity.

10. What other coined economies teach

An “economy” label survives when it names a scarce resource, actor, transaction, or value shift more clearly than competing language.

Label	Core unit or mechanism	Why it held or spread	Failure or caution	Lesson for Operator Economy
Attention economy [K1]	Abundant information makes attention scarce	Clear scarcity logic; applies across media and advertising	Broad enough to become descriptive; weak source exclusivity	Name the scarce control point, not merely the participant
Experience economy [K2]	Firms stage memorable experiences as a distinct economic offering	Actionable value progression and clear managerial consequence	Became broad consulting language	Give firms a decision model, not only a cultural observation
Platform or gig economy [K3]	Digitally mediated matching of labor, assets, and demand	Measurable transactions, firms, workers, and policy consequences	“Sharing” rhetoric often hid rental and labor control	Define who sets price, owns demand, and bears risk
Creator economy [K4]	Individuals monetize content, audience, and influence through platforms	Clear actor and visible tools, platform maps, and payout data	Catchall boundaries; platform dependence undercut autonomy	Do not equate independence with account ownership
Passion economy [K5]	Platforms let people monetize differentiated skills and individuality	Strong contrast with standardized gig work; practical platform features	Partly absorbed into creator/solo-work language	Show the institutional mechanism and metrics, not just identity
Ownership economy [K6]	Users who build and fund networks also own them through crypto mechanisms	Specific value-distribution thesis	Technology-cycle dependence and unstable definitions; its originator later narrowed the claim	Ownership must specify actual rights, not vibes
Subscription economy [K7]	Firms shift from product sales to recurring service relationships	Product infrastructure, book, institute, events, and a recurring index built from data for more than 600 companies	Category adoption required years of institution-building and maintained use	The best comparator: product + data + repeated research + source identity
Agentic economy [K8]	Consumer and business agents transact programmatically	Clear proposed transaction and governance questions	Emerging, autonomy and accountability unsettled	Keep human accountability as an explicit boundary
Knowledge economy [K9]	Production, distribution, and use of knowledge	Government taxonomies, indicators, and long-running policy institutions	So broad that boundaries and pricing remain difficult	Experienced judgment is an input, not a new Operator-Economy transaction by itself
Company of One / Small Giants / Zebras [K10]	Right-sized, enduring, or stakeholder-oriented firm archetypes	Memorable identity plus books, movements, cases, and practices	They are firm philosophies, not measured economies	A distinct movement or method may be more ownable than an “economy” noun
Hidden Champions [K11]	Low-visibility firms dominate narrow global niches	Long-running cases, measurable market position, strategic mechanisms	Usually established employers, not tiny operator cores	“Too small to notice” has deep prior art; narrowness alone is not new

Trademark status is not category durability

Federal records reinforce the distinction between a durable idea and a maintained source mark. These are live status snapshots, not holdings that any phrase is legally protectable or unprotectable:

Exact phrase record	Live federal status on 24 August 2026	What it actually shows
EXPERIENCE ECONOMY, Serial 85878999 [M1]	Registration cancelled because an acceptable Section 8 declaration was not successfully filed	Required maintenance must be completed; this record does not establish whether marketplace use continued
PASSION ECONOMY, Serial 98712811 [M2]	Application abandoned after failure to respond	Popularity of a category phrase does not complete prosecution
THE OWNERSHIP ECONOMY, Serial 88918714 [M3]	Application abandoned after no Statement of Use	A thesis and a bona fide source-use specimen are different things
CREATOR ECONOMY, Serials 90140724 and 85583277 [M4]	Exact-phrase applications are dead; composite marks may differ	A widely used category can spread while exact-phrase exclusivity weakens
SUBSCRIPTION ECONOMY, Serial 85253216 [M5]	Active and renewed in Class 35; Class 36 and 38 portions cancelled; first use claimed in 2009	The credible counterexample paired early use with products, an institute, events, and recurring platform data

Closest conceptual prior art

The organizational form is old. Charles Handy's 1989 **Shamrock Organization** described a professional core, contractual fringe, and flexible labor force [P1]. Walter Powell's 1990 network-form work analyzed exchange across boundaries that were neither clean market nor hierarchy [P2]. The virtual-corporation literature described firms coordinating external specialists decades before generative AI [P3].

The closest recent commercial collision is the 2021 argument that **“the software stack is the new firm.”** It already describes experienced professionals going solo, software replacing operational support, platforms supplying demand and networks, vertical stacks acting as systems of record, and customer acquisition as the hardest problem [P4].

The narrower contribution this paper can attempt is:

- it maps the entire capability perimeter, including AI, people, capital, logistics, and regulated institutions;
- it separates capacity from accountability;
- it measures the independence lost through rented infrastructure;
- it treats key-person concentration as a liability;
- it distinguishes sufficiency from moat;
- it publishes a ledger and falsifiers rather than a platform investment thesis.

That is a researchable difference. It is not yet a demonstrated causal contribution, and it is not a claim to have discovered solo work, modular firms, networked firms, or small-company ambition. The residual contribution may prove to be a dependency, accountability, and risk-incidence audit method rather than a new economic category.

The seven-part durability test

A category holds when it develops:

- 1 a hard exclusion boundary;
- 2 a distinct actor or unit;
- 3 a causal mechanism;
- 4 measurable proxies and a denominator;
- 5 repeated external or platform data;
- 6 products, institutions, or an ecosystem that use the concept;
- 7 a public correction cadence.

The current Operator Economy has an audience identity and media property. It does not yet have the denominator, external adoption, longitudinal data, or correction institution.

11. Exact-name and semantic collisions

The current public field includes at least:

Use	Meaning	Timing evidence
CoinDesk / Ivo Entchev	Web3 upgrade to the gig or sharing economy; operators supply labor, hardware, data, compute, and storage	Published 26 April 2024 [N1]
Intellivance / Danny Matulula	AI-enabled fractional operators serving three to five clients	Published 17 December 2025 [N2]
Patrick J. Murphy	Execution, alignment, and leadership as value creation	Current page; priority date not established here [N3]
The Operator Class	A system-led commerce and revenue category with manifesto, archetypes, index, calculator, and diagnostic	Current 2026 site [N4]
Ivanooo / Direction Economy URL lineage	AI direction capacity, a named index, formula, methodology, and replication logic at an earlier AI-operator URL	Current page; reliable publication priority not established here [N5]
Ralph Quintero	One person plus an AI-agent fleet appears like a larger agency; creator economy shifts toward operator economy	LinkedIn post dated 16 March 2026 [N6]
bstarOne	“AI Operator Economy” and Vanguard framework for agencies deploying AI business systems	Press release/current commercial site dated 17 March 2026 [N7]

The phrase therefore has at least three materially different meanings:

- 1 network or gig participants;
- 2 AI-enabled fractional labor;
- 3 accountable business architects or owner-operators.

That fragmentation is strategically bad for a broad exclusivity claim. It may still be usable as a media identity, but it is not honest to present it as a newly coined phrase.

12. What is actually distinctive in this draft


Established lineage — do not claim

- transaction-cost and firm-boundary theory;
- outsourced and networked firms;
- cloud and software lowering fixed commitments;
- one-person and right-sized companies;
- vertical SaaS and systems of record;
- small niche incumbency;
- distribution, trust, switching costs, and data context;
- venture power laws and alternative scoreboards;
- AI augmentation and workflow integration.

Distinctive synthesis — claim carefully

- 1 **Accountability-capacity mapping:** the core retains customer promise and recourse while duties, liability, and capacity are mapped across the perimeter.
- 2 **A moderated boundary-shift hypothesis:** capability granularity may permit a smaller permanent boundary, but asset specificity, integration, learning, security, regulation, and bargaining power can reverse the move.
- 3 **Dependency-Adjusted Independence assessment:** ownership independence is examined against operational revocability, concentration, recovery time, and tested substitution.
- 4 **Sufficiency-versus-defense separation:** normalized economics, reachable demand, and chosen delivery capacity are tested separately from entrant threats.
- 5 **The five-ledger method:** economics, control, resilience, accumulated position, and labor/risk incidence are evaluated without premature composite scoring.
- 6 **A disprovable category contract:** the classifier, prevalence threshold, matched comparisons, and correction rules can reject the economy claim.

These elements may become distinctive as an integrated assessment and expression. They are not yet validated contributions, and that is different from being unprecedented, patentable, or trademark-exclusive.



13. Predictions and research program

The thesis earns credibility only if it makes risky predictions. The table below is a study-design inventory, not a preregistration. Before collecting outcome data, every test must fix the sample, exposure, comparison group, time horizon, minimum effect, uncertainty rule, and rejection threshold.

Prediction	Measurement	What would weaken the thesis
High-modularizability sectors show a larger post-2022 rise in eligible operator firm-offer-year units than matched low-modularizability sectors	Repeated classified samples, formation, payroll, nonemployer, contractor/vendor, and survival data	No differential change after controls, or classification is not reproducible
Quality-adjusted value added improves relative to total labor-hours, total input dollars, and owner-supplied capital—not merely payroll headcount	Output quality, revenue and intermediate inputs, payroll, allocated supplier labor, contractor/model/tool spend, owner hours, and capital	Apparent productivity disappears when all inputs and quality are counted
Broad workflow integration predicts performance better than casual tool adoption	Number and depth of integrated functions, controls, outcomes	Tool access alone explains durable performance
Direct, portable buyer access predicts survival and margin	Channel concentration, customer identity portability, CAC/payback, retention	Platform-dependent firms perform equally after risk controls
Dependency concentration predicts outages, margin shocks, or firm failure	Vendor/platform concentration, price changes, revocations, replacement time	Concentration has no material outcome effect
Operator key-person exposure predicts lower resilience and transfer value	Critical decision coverage, absence tests, sale multiples, interruption data	Documentation and delegation show no relationship
Embedded workflow position predicts retention only when customer value remains positive	Integration breadth, switching cost, NPS/complaints, retention	Lock-in without value performs equally or better over time
Small market size deters only some scaled de novo entrants	Entry by peer, adjacent, platform, customer, and acquirer class	"Too small" consistently deters all relevant entrant types

The economy claim fails for this research program if any of these persist after a powered, preregistered study:

- 1 the classifier cannot reach 0.80 agreement across independent coders;
- 2 estimated prevalence does not clear the promotion threshold or does not rise across repeated periods;
- 3 high- and low-modularizability sectors show no material differential boundary change;
- 4 the productivity result disappears after allocating supplier labor, owner labor, capital, and quality;
- 5 matched conventional small firms show the same control, dependency, incidence, and outcome profile.

Minimum viable research design

Dataset

- SUSB employer firms and establishments;
- Nonemployer Statistics, with receipts-cutoff and NAICS-vintage warnings;
- Business Dynamics Statistics for entry, exit, and age;
- Annual Business Survey and BTOS for technology use and owner characteristics;
- occupational and contractor/vendor inputs where available;
- original firm-offer-year boundary ledgers.

These public datasets cannot observe the defining construct: customer-promise control, residual decision rights, supplier capacity, portability, mapped accountability, or owner labor. They can frame the outer population and comparisons; they cannot produce an Operator Economy denominator by themselves.

Do not casually add the populations together. Define the firm $\times$ offer $\times$ year unit, industry vintage, coding window, shared-supplier rule, and exclusions before calculating a headline.

Field work

Use four stages:

- 1 **Taxonomy discovery:** 30 structured interviews across the five configurations below, with at least one-third failed, dormant, rejected, or counter cases. This can refine definitions; it cannot estimate prevalence.
- 2 **Reliability pilot:** 100 cases coded independently by at least two people; revise the manual until agreement reaches 0.80 without adjudication gaming.
- 3 **Outcome panel:** at least 300 operator and matched non-operator firm-offer-year units observed quarterly for 24 months. Power analysis, not the desired headline, sets the final sample.
- 4 **Denominator study:** a probability-based survey or a data partnership with a processor, platform, insurer, payroll provider, or government agency. Until this exists, make no national size claim.

The five discovery configurations are:

- 1 independent professional with a vertical stack;
- 2 software or digital-asset operator;
- 3 productized service or small agency;
- 4 local or physical operator using rented logistics and software;
- 5 regulated operator using licensed institutional partners.

For every firm, map:

- permanent core;
- total human and institutional perimeter;
- AI and software functions;
- demand source;
- customer and cash control;
- critical dependencies;
- operator hours and decisions;
- total internal and allocated supplier labor-hours;
- normalized economic surplus and owner-supplied capital;
- accountability and liability by layer;

- contractor pay, benefits, volatility, classification, and training burden;
- failures, substitutions, and recovery time.

Counter-case policy

At least one-third of published cases should challenge the thesis:

- AI adoption that increased complexity;
- platform removal or repricing;
- cheap clone with better distribution;
- small firm that was not economically sufficient;
- operator burnout or key-person failure;
- customer insourcing;
- large incumbent that integrated AI more effectively;
- regulation that forced internal capability or scale.

If counter-cases are treated as exceptions to be explained away, this becomes ideology.

14. Protection architecture

What can and cannot be protected

The economic idea itself cannot be made exclusive by trademark. Trademark protects a source identifier used for goods or services [T2]. Copyright may protect human-authored wording and graphics, human modifications, and sufficiently creative human selection or arrangement; it does not protect the underlying idea, system, method, facts, or wholly AI-generated material [T5; T8].

A live exact-word USPTO search on 24 August 2026 returned no federal record for “Operator Economy” or “The Operator Economy” [T1]. That is not clearance. Federal search does not erase common-law use, commercial websites, publications, state records, international rights, or confusingly similar marks [T6].

The phrase faces three practical problems:

- 1
- it directly names the subject matter of the publication and may be viewed as descriptive [T3];
- 2
- varied third-party use can make it informational rather than source-identifying [T3];
- 3
- adding “The” generally does little to create distinctiveness [T4].

The title of one PDF or book is also generally not registrable as the mark for that single work. A genuine recurring series can function differently when consumers recognize the title as a source [T7].

A more credible architecture has separate layers:

Layer	Role	Strategic treatment
Operator Economy	Media identity and open category hypothesis	Use if strategically valuable, but do not claim the phrase was coined here or assume broad exclusivity
Distinctive method or series mark	Identifies recurring education, assessment, or research services	Coin only after naming work and full legal clearance; avoid generic “Operator Index” and “Operator Diagnostic” collisions
Human-authored expression	Working papers, diagrams, taxonomy, case narratives, course and publication text	Document human authorship, modifications, and creative selection or arrangement; preserve dated source files and versions
Proprietary dataset	Longitudinal boundary, dependency, economics, and outcome records	This is the strongest practical moat; publish enough method for credibility while protecting private respondent data
Calibrated instrument	Repeatable ledger or benchmark based on observed outcomes	Do not invent weights before calibration; transparent method beats pseudo-science
Visual source identity	A consistent, distinctive, nonfunctional appearance	Evaluate trade-dress distinctiveness and nonfunctionality separately; visual consistency alone does not create protection
Provenance archive	Versions, hashes, corrections, authorship records, and use specimens	Maintain as evidence and governance controls, not as trade dress
Trade secret	Nonpublic scoring weights, research operations, or customer-specific implementation	Use only where secrecy creates value; public doctrine cannot also be a secret

Current risk map

Asset	Present assessment	Reason
"The Operator Economy" for broad educational/media exclusivity	High risk / low confidence	Exact phrase has earlier and current third-party use; meaning is fragmented; the words describe the topic
A single thesis PDF title	Poor trademark vehicle	Single-work title rule and weak source-identifying use
A recurring publication or service under a cleared mark	Possible, fact-dependent	Requires distinctiveness, actual or bona fide intended goods/services, consistent use, and clearance
Human-authored prose, diagrams, modifications, and creative arrangement	Potentially protectable expression	Wholly AI-generated elements, ideas, methods, names, titles, and short phrases are outside that claim [T5; T8]
A concrete technical invention	Separate patent analysis only	A business thesis or abstract framework is not made patentable by adding software language
Confidential case data and weights	Potential trade-secret value	Only while valuable information is actually kept secret with reasonable controls

Naming rule

The category term and the owned mark should not be forced to do the same job.

The category should be understandable and shareable. The owned mark should be distinctive enough to identify your specific method or series. A coined or suggestive sub-brand may have better legal prospects than another descriptive "Operator + generic noun" name.

TM or SM may be used before filing, but using one here would assert a source-mark claim before clearance and a deliberate brand decision. Reserve ® for a federal registration and only for the goods or services listed in that registration [T2]. A preliminary web search is not legal clearance.

Discard the obvious candidate set. Current businesses already use or promote:

- The Operator Class;
- Operator Index;
- Operator Score and Operator Scorecard;
- Operator Manifesto;
- Operator Diagnostic;
- Operator Library;
- five Operator Archetypes;
- Operator Operating System or equivalent system language.

The proprietary name should not be another "Operator + generic noun" combination.

Construct-name collision check

The first naming pass also killed two labels created inside this paper:

Working label	Current evidence	Treatment
Sufficiency Band	Zero exact live federal word-mark records, but the exact term already exists in plant-nutrition diagnostics; active SUFFICIENCY CURVE registration in financial-sufficiency services sits in the same natural naming family [C1]	Retired as a proprietary-looking name; use the descriptive “viability gate” for the break-even check
Operator Boundary Model	Zero exact live federal word-mark records, but “Operator Boundary” already appears in an AI-governance standard [C2]	Retired as a claimed coined name; use “boundary-ledger method” descriptively
Dependency-Adjusted Independence	Zero exact live federal word-mark records and no exact indexed full-phrase use found in this pass; “dependency-adjusted” is an established analytic modifier [C4]	Keep as the strongest provisional technical term, not as a cleared trademark
Small Core, Thick Perimeter	No exact record or indexed exact use found, but materially similar small-core/network models and close small-core rhetoric exist [C3]	Keep only as an unclaimed working tagline pending similar-mark and common-law clearance

An exact-record search is a narrow screen. It does not clear similar marks, unregistered use, state or foreign rights, domains, courses, publications, or failure-to-function risk.

Filing sequence

- 1 Define the recurring goods or services that the mark will identify.
- 2 Run a professional clearance search across federal, state, common-law, domain, app, publication, course, and international uses relevant to the launch.
- 3 Select a distinctive mark; avoid names already used for operator indexes, diagnostics, classes, or manifestos.
- 4 Preserve first-use evidence and consistent specimens.
- 5 File only after counsel assesses distinctiveness, ownership, classes, and collision risk.
- 6 Keep the category language honest even if a registration issues; registration is not proof of authorship of the idea.

Possible U.S. classes depend on actual facts, not aspiration: Class 41 may cover defined education or non-downloadable publication services; Class 9 may cover downloadable publications; Class 16 may cover print; Class 35 should be claimed only for actual business or consulting services, not merely content about business.

This is research guidance, not legal advice.

15. The prioritized plan

Priority 0 — Stop the false claims now

- Do not say “completely unique,” “coined here,” “first,” or “trademarked thesis.”
- Do not reuse “most software companies have four people or fewer.”
- Do not call the SUSB firms solo, typical, durable, or AI-created.
- Do not call smallness a moat.
- Do not publish the retained 38-page reference PDF.

Priority 1 — Repair the evidence and provenance in 48 hours

- recover the exact live document post and clean attachment;
- capture lifetime performance;
- hash and archive the public bytes;
- correct the false publication pointer;
- create a working-paper version ledger;
- preserve this paper as Research Draft 0.1, not canonical doctrine.

Priority 2 — Freeze the category contract in one week

Approve or reject:

- working definition of operator firm;
- inclusion test and exclusions;
- human accountability boundary;
- category-versus-media-brand distinction;
- the seven propositions and their falsifiers;
- replacement positioning sentence.

Proposed public sentence:

The Operator Economy is a research program on firms where an owner-controlled core of no more than ten people retains customer recourse while assembling material external capability—and measures the control, labor, liability, and risk that remain outside that core.

Priority 3 — Build the taxonomy over six weeks

- 30 structured operator interviews;
- five configuration groups;
- at least 30 complete five-ledger cases;
- one-third counter-cases;
- documented owner and supplier labor, liability, capital, and dependency failures;
- no anonymized “illustrative” case presented as observed fact.

Priority 4 — Prove the classifier over the next eight weeks

- publish the unit-of-analysis memo first;
- construct reproducible SUSB and NES extracts;
- separate employer, nonemployer, establishment, firm, and worker;
- code at least 100 cases independently;
- require 0.80 agreement before outcome analysis;
- freeze thresholds, exclusions, and shared-supplier rules;
- release the coding manual, disagreements, calculations, and caveats.

Priority 5 — Build the evidence program over 18–24 months

- recruit at least 300 operator and matched non-operator firm-offer-year units, subject to power analysis;
- observe them quarterly for 24 months, including exits and dormant firms;
- secure a probability survey or provider/government data partnership before making a national denominator claim;
- preregister effect sizes, timing, and rejection rules;
- run the five-ledger method and track whether the measures predict:
 - normalized economic surplus;
 - retention;
 - interruption;
 - operator hours;
 - replacement time;
 - labor and risk incidence;
 - ability to transfer or sell the firm.

Only then consider a composite benchmark. An uncalibrated index would repeat the exact mistake this thesis is trying to correct: visual authority standing in for evidence.

Priority 6 — Earn the mark after the substance

After recurring commercial or educational use exists:

- commission full clearance;
- choose a distinctive method or series mark;
- register the expression, archive, and specimens correctly;
- keep “Operator Economy” as an open category/media term unless counsel finds a narrower defensible claim.

16. Final strategic judgment

The original paper's emotional conclusion is worth keeping:

A business can be the right size for the life, market, and responsibility it was built to support.

But that is a normative choice, not economic proof.

The full thesis is harder:

Institutional capability is becoming divisible and rentable in some work. Under favorable conditions, a small permanent core may therefore coordinate a productive system much larger than its payroll. The form is old; the proposed change is wider diffusion and a lower feasible core size in compatible sectors. Fixed commitments may fall while coordination, verification, platform, vendor, labor-incidence, regulatory, and key-person risks are added or reconfigured. Production access can become less scarce, but market power, rights, contracts, regulation, trust, judgment, workflow position, and accountability determine who captures value. A viable operator firm must pass normalized economics, reachable-demand, and chosen-capacity tests, then survive a separate entrant and dependency screen. Its durability must be measured; being too small to notice is not a moat.

That hypothesis is now specific enough to enter serious scrutiny. It has not survived it yet.

The next level is not a better slogan. It is a body of evidence that other people can use, test, criticize, and cite without needing to trust the founder.

Appendix A — Repository evidence

- operator-economy/studio/originate/too-small-to-bother/reference/working-paper-01-print-capture.pdf
- operator-economy/studio/originate/too-small-to-bother/research.md
- [R1] operator-economy/research/briefs/ep005-operator-economy-thesis.md, legacy SUSB calculation; universe meaning cross-checked against [E1–E2], raw extract not rebuilt in this pass
- brown-man-content/content-packages/week-of-2026-08-03/oe-launch/2-feed-post-monday-1230.md
- brown-man-content/content-packages/week-of-2026-08-10/monday/article.md
- brown-man-content/channels/linkedin/tracking/opener-tracker.csv, rows 115, 125, and 127
- brown-man-content git object 4552957:content-packages/week-of-2026-08-10/monday/post.md
- content-os/facts.md, Operator-economy thesis research section

Appendix B — Claim-level external source notes

Bracketed IDs map directly to claims in the paper. Locators identify the relevant page, section, abstract, record, or visible page element. All web sources were accessed 24 August 2026. “Current page” evidence establishes present use only; without an archived copy or exposed publication date, it does not establish priority. Confidence describes the narrow claim made here, not the source as a whole.

Firm boundaries, small cores, and AI evidence

- [E1] U.S. Census Bureau, *Statistics of U.S. Businesses Methodology*. Locator: “Business Register Source” and “Scope and Coverage”; the register contains establishments with paid employees and excludes businesses without employees. Confidence: high for the statistical universe, not the legacy paper’s calculation. www.census.gov/programs-surveys/susb/technical-documentation/methodology.html
- [E2] U.S. Census Bureau, *Nonemployer Statistics Methodology*. Locator: “Nonemployer Universe,” “Maximum and Minimum Receipts Cutoffs,” and “Industry Classification”; NES uses administrative tax records for businesses with no paid employees and applies changing cutoffs and NAICS vintages. Confidence: high for the separate universe and comparability warning. www.census.gov/programs-surveys/nonemployer-statistics/technical-documentation/methodology.html
- [E3] Grundy, Breaux, and Khatiwoda, U.S. Census Bureau, *Large Firms With at Least 20 Employees Biggest AI Users*, 26 May 2026. Locator: “U.S. Business AI Use” and Figure 2; overall use hovered between 17% and 20%, use changed significantly only for firms with at least 20 employees, and fewer than 20% of firms with four or fewer employees reported use. Confidence: high, descriptive BTOS result. www.census.gov/library/stories/2026/05/ai-use-businesses.html
- [E4] Bonney et al., *The Microstructure of AI Diffusion*, U.S. Census CES Working Paper 26-25, April 2026. Locator: page abstract; 18% of firms used AI, 66% of users reported augmentation only, AI-related employment decreases occurred in 2%, and integration breadth correlated with commercial performance. Confidence: high for the survey report; correlation is not causation. www.census.gov/library/working-papers/2026/adrm/CES-WP-26-25.html
- [E5] OECD, *Artificial Intelligence and Competitive Dynamics in Downstream Markets*, 2025/2026 release. Locator: Chapter 2 discussion of adoption by firm size, economies of scale, data and compute, integration, lock-in, and entry. The 39%-versus-12% comparison produces the reported 3.3-fold gap. Confidence: moderate to high synthesis; cross-country area measures are not one universal firm census. www.oecd.org/en/publications/artificial-intelligence-and-competitive-dynamics-in-downstream-markets_ccf0624a-en/full-report/component-5.html
- [E6] Marchesi and Tang, *When AI Meets Entrepreneurship*, working paper revised 9 June 2026. Locator: abstract and difference-in-differences result reporting more than 10% higher firm formation in AI-compatible industries after ChatGPT.

Confidence: provisional; this is a working paper and sector-compatibility design, not settled economy-wide causation.
papers.ssrn.com/sol3/papers.cfm?abstract_id=5759264

- **[E7]** Forman and van Zeebroeck, *Firm Size and Boundaries*, NBER chapter. Locator: PDF pages 5 and 11; cloud converts some data-center fixed cost into usage-based cost, while off-the-shelf cloud services can lower minimum efficient scale. Confidence: high for lineage and conditional mechanism.
www.nber.org/system/files/chapters/c14017/revisions/c14017.rev0.pdf
- **[P1]** Charles Handy, *The Age of Unreason*, 1989. Locator: the Shamrock Organization chapter describing a professional core, contractual fringe, and flexible labor force; pagination varies by edition. Confidence: moderate because the linked record is bibliographic, not a full-text edition. books.google.com/books/about/The_Age_of_Unreason.html?id=TQhPAAAAMAAJ
- **[P2]** Walter W. Powell, *Neither Market Nor Hierarchy: Network Forms of Organization*, 1990. Locator: pp. 295–336, especially the opening definition and comparison of network, market, and hierarchy forms. Confidence: high for organizational lineage.
web.stanford.edu/~woodyp/powell_neither.pdf
- **[P3]** William H. Davidow and Michael S. Malone, *The Virtual Corporation*, 1992. Locator: book-level thesis; linked record is bibliographic and pagination is edition-specific. Confidence: moderate for lineage, not empirical prevalence.
books.google.com/books/about/The_Virtual_Corporation.html?id=tvD2xcq34JIC
- **[P4]** Amble, Coolican, and Rampell, *As More Workers Go Solo, the Software Stack Is the New Firm*, 24 September 2021. Locator: opening thesis, “A New Stack Emerges,” and discussion of demand aggregation and customer acquisition. Confidence: high for a close published commercial thesis collision.
a16z.com/as-more-workers-go-solo-the-software-stack-is-the-new-firm/

Coined economies and adjacent movements

- **[K1]** Herbert Simon, *Designing Organizations for an Information-Rich World*, 1971, discussion of attention scarcity; Michael Goldhaber, *The Attention Economy and the Net*, 1997, opening definition. Confidence: high for conceptual lineage.
gwern.net/doc/design/1971-simon.pdf and firstmonday.org/ojs/index.php/fm/article/view/519
- **[K2]** Pine and Gilmore, *Welcome to the Experience Economy*, Harvard Business Review, July–August 1998. Locator: the commodity-to-goods-to-services-to-experiences progression and staging principles. Confidence: high for category mechanism and management model. hbr.org/1998/07/welcome-to-the-experience-economy
- **[K3]** OECD, online-platform measurement roadmap. Locator: Chapter 7.4, definitions and proposed measurement of platforms, workers, producers, and transactions. Confidence: high for the need for countable actors and flows.
www.oecd.org/en/publications/measuring-the-digital-transformation_9789264311992-en/full-report/component-62.html
- **[K4]** SignalFire, *Creator Economy Market Map*, 2024, maps participant and company classes; Stripe's 2021 and 2023 analyses use payout data from creator platforms. Confidence: high for external ecosystem mapping and payment measurement, not the whole creator universe. www.signalfire.com/blog/creator-economy ; stripe.com/blog/creator-economy ; stripe.com/blog/creator-economy-2023
- **[K5]** Li Jin, *The Passion Economy and the Future of Work*, 8 October 2019. Locator: contrast with standardized gig work and five platform-design features. Confidence: high for the original category framing.
a16z.com/the-passion-economy-and-the-future-of-work/
- **[K6]** Jesse Walden, *The Ownership Economy*, 14 July 2020, defines users building, operating, funding, and owning networks; *Everything Is Market*, 13 February 2026, says ownership materialized in significant but narrow financial verticals and revises the thesis. Confidence: high for original claim and authorial correction.
variant.fund/articles/the-ownership-economy-crypto-and-consumer-software/ and variant.fund/articles/everything-is-market/
- **[K7]** Zuora, *Subscription Economy Index* and *Subscribed Institute*. Locator: current index methodology uses aggregated data from more than 600 companies; the institute is the recurring research institution. Confidence: high for product-plus-data institutionalization. www.zuora.com/resource/subscription-economy-index/ and www.zuora.com/about/subscribed-institute/
- **[K8]** Microsoft Research, *The Agentic Economy*, May 2025. Locator: abstract and paper's transaction/governance framing. Confidence: high for the proposed unit; low for mature prevalence because the category is emerging.
www.microsoft.com/en-us/research/publication/the-agentic-economy/

- **[K9]** Fritz Machlup, *The Production and Distribution of Knowledge in the United States*; OECD, *The Knowledge-Based Economy*, 1996. Locator: book-level production/distribution taxonomy and OECD measurement framework. Confidence: high for established lineage. assets.press.princeton.edu/about_pup/PUP100/book/fulltext.pdf and [one.oecd.org/document/OCDE/GD\(96\)102/En/pdf](https://one.oecd.org/document/OCDE/GD(96)102/En/pdf)
- **[K10]** Brandel, Zepeda, Scholz, and Williams, *Zebras Fix What Unicorns Break*, 2017; Paul Jarvis, *Company of One*; Bo Burlingham, *Small Giants*. Locator: movement manifesto and book-level firm philosophies. Confidence: high that right-sized, enduring, non-unicorn firm identity has extensive prior art. resources.platform.coop/en/resources/zebras-fix-what-unicorns-break/; www.penguinrandomhouse.com/books/566626/company-of-one-by-paul-jarvis/; www.penguinrandomhouse.com/books/18048/small-giants-by-bo-burlingham/
- **[K11]** Hermann Simon, *Hidden Champions*, 1996. Locator: book-level case program on low-visibility niche market leaders. Confidence: high for the prior-art category; the linked record is bibliographic. books.google.com/books/about/Hidden_Champions.html?id=OrfLkELV34C

Federal status snapshots for comparator phrases

- **[M1]** EXPERIENCE ECONOMY, U.S. Serial 85878999. Locator: TSDR status and maintenance record; registration cancelled because an acceptable Section 8 declaration was not successfully filed. The record does not establish whether marketplace use continued. Confidence: high for status on access date, not phrase-level protectability. tsdr.uspto.gov/
- **[M2]** PASSION ECONOMY, U.S. Serial 98712811. Locator: TSDR status; application abandoned after failure to respond. Confidence: high for status on access date. tsdr.uspto.gov/
- **[M3]** THE OWNERSHIP ECONOMY, U.S. Serial 88918714. Locator: TSDR status; application abandoned for no Statement of Use. Confidence: high for status on access date. tsdr.uspto.gov/
- **[M4]** CREATOR ECONOMY, U.S. Serials 90140724 and 85583277. Locator: TSDR status pages; both exact-phrase records are dead. Confidence: high for these records only; composites and other owners require separate searches. tsdr.uspto.gov/ and tsdr.uspto.gov/
- **[M5]** SUBSCRIPTION ECONOMY, U.S. Serial 85253216. Locator: TSDR status and registration maintenance; active/renewed in Class 35, with the Class 36 and Class 38 portions cancelled, and first use claimed in 2009. Confidence: high for this federal record, not a claim to ownership of the phrase in every context. tsdr.uspto.gov/

Exact phrase and current operator-category uses

- **[N1]** Ivo Entchev, *DePIN for the Win*, CoinDesk, 26 April 2024. Locator: paragraph beginning “what I will call the ‘operator economy’,” defining it as what is more commonly called the gig or sharing economy, plus later “Web3 operator economy” uses. Confidence: high for public exact-phrase use by that date; editorial use alone is not trademark priority. www.coindesk.com/opinion/2024/04/26/depin-for-the-win-spreading-the-benefits-of-the-gig-economy
- **[N2]** Danny Matulula, *The Operator Economy: Why the Cost of Labor Is About to Collapse*, 17 December 2025. Locator: opening definition and fractional-operator example of serving three to five clients. Confidence: high for the dated published meaning. intellivance.ai/blog/the-operator-economy
- **[N3]** Patrick J. Murphy, *The Operator Economy*. Locator: current commercial landing page and operator-leadership services. Confidence: high for present commercial use; priority date not established from this page. www.patrickjmurphy.com/operator-economy
- **[N4]** *The Operator Class*. Locator: current home and library pages using “The Operator Economy,” Operator Index, Operator Score, Operator Manifesto, Operator Diagnostic, Operator Library, and five archetypes. Confidence: high for present strategic collision; exact publication priority is not established here. theoperatorclass.com/
- **[N5]** Ivanooo, current page at the former AI-operator URL and methodology page. Locator: named index, two-axis construct, formulas, measured-versus-modelled ledger, and replication instructions. Confidence: high for present conceptual collision; current page title and URL lineage do not establish its original publication date. ivanooo.com/ai-operator-economy and ivanooo.com/ai-operator-economy/methodology

- **[N6]** Ralph Quintero, LinkedIn post dated 16 March 2026. Locator: post body describing one person with an AI-agent fleet appearing like a larger agency and “creator economy” shifting toward “operator economy.” Confidence: high for dated current-platform use; platform availability may change. www.linkedin.com/posts/ralphquintero_i-met-someone-last-week-running-what-looked-activity-7439294987814612992-f9rM
- **[N7]** bstarOne, current commercial site and dated 17 March 2026 release material. Locator: “AI Operator Economy” and Vanguard framework language for agencies deploying AI business systems. Confidence: high for present commercial use. bstarone.com/

Legal and naming guidance

- **[T1]** USPTO Trademark Search, exact live word searches for OPERATOR ECONOMY and THE OPERATOR ECONOMY run 24 August 2026; zero exact federal records returned. No static result was retained. Confidence: high for that narrow live query only; this is not clearance. tmsearch.uspto.gov/
- **[T2]** USPTO, *What is a trademark?* Locator: definition of a trademark as identifying goods or services and their source; TM and SM may be used without filing, while ® is reserved for a federally registered mark and the covered goods or services. Confidence: high, primary agency guidance. www.uspto.gov/trademarks/basics/what-trademark
- **[T3]** USPTO, *Strong trademarks*, and TMEP 1202.04, *Informational Matter*. Locator: descriptive-strength discussion and refusal where wording fails to function as a mark because consumers perceive information rather than source. Confidence: high, primary agency guidance. www.uspto.gov/trademarks/basics/strong-trademarks and tmept.uspto.gov/RDMS/TMEP/print?href=TMEP-1200d1e1927.html&version=current
- **[T4]** Trademark Trial and Appeal Board, Serial 87809270, decision p. 3. Locator: the definite article “THE” generally adds no source-indicating significance to otherwise descriptive wording. Confidence: high for the stated general principle; facts differ by application. tab-reading-room.uspto.gov/cms/rest/legal-proceeding/87809270/decision/EXA_13.pdf
- **[T5]** U.S. Copyright Office, Circular 33, *Works Not Protected by Copyright*. Locator: ideas, methods, systems, names, titles, and short phrases. Confidence: high, primary agency guidance. www.copyright.gov/circls/circ33.pdf
- **[T8]** U.S. Copyright Office, *Copyright and Artificial Intelligence, Part 2: Copyrightability*, January 2025. Locator: pp. ii–iii and Parts II–III; copyright can cover human-authored expression, sufficiently creative human selection or arrangement, and human modifications, but not wholly AI-generated material, and prompting alone does not provide sufficient control. Confidence: high, primary agency report; application remains fact-specific. www.copyright.gov/ai/Copyright-and-Artificial-Intelligence-Part-2-Copyrightability-Report.pdf
- **[T6]** USPTO, *Federal trademark searching*. Locator: search scope and need to investigate beyond exact federal wording. Confidence: high, primary agency guidance. www.uspto.gov/trademarks/search/federal-trademark-searching
- **[T7]** USPTO, *Title of a Single Work Refusal and How to Overcome This Refusal*. Locator: single-work title rule and series distinction. Confidence: high, primary agency guidance. www.uspto.gov/trademarks/laws/title-single-work-refusal-and-how-overcome-refusal
- **[C1]** SUFFICIENCY CURVE, U.S. Serial 87072578, active registration for financial-sufficiency services; Sousa et al., 2021, uses “sufficiency bands method” in plant-nutrition diagnosis. Locator: TSDR status and article title/method. Confidence: high that the naming family is pre-existing; legal similarity requires counsel. tsdr.uspto.gov/ and rsdjournal.org/index.php/rsd/article/download/14805/13370/194326
- **[C2]** ARCS v1.0, ARCS-OPB. Locator: “Operator Boundary” control in the AI-governance standard. Confidence: high for same-field terminology collision, not trademark rights. arcsstandard.org/downloads/ARCS_v1_0_Full_Standard.pdf
- **[C3]** Stripe, *Small company structure explained*, section on network/lean structures; Traversally, current Team Operating Model. Locator: small core with outsourced functions and current “Small core. Broad capability. Clear accountability.” wording. Confidence: high that the organizational idea and close small-core rhetoric are not unique. stripe.com/resources/more/small-company-structure-explained and traversally.ai/team-operating-model
- **[C4]** Romano and Wolf, 2015 working paper, uses dependency-adjusted p-values and thresholds. Locator: abstract and method. Confidence: moderate that “dependency-adjusted” is established analytic wording; the source does not use the full “Dependency-Adjusted Independence” phrase or establish trademark rights. ideas.repec.org/p/jgu/wpaper/1505.html

Additional legal boundaries

- USPTO, Goods and services: www.uspto.gov/trademarks/basics/goods-and-services
- USPTO MPEP 2106, Patent Subject Matter Eligibility: www.uspto.gov/web/offices/pac/mpep/s2106.html
- USPTO, Trade Secret Policy: www.uspto.gov/ip-policy/trade-secret-policy

Appendix C — Research limits

- This is a U.S.-first empirical and legal review with global conceptual prior art.
- The exact live 10 August document post and clean attachment were unavailable.
- Several current web pages do not expose reliable first-publication dates.
- Search results are not a comprehensive common-law or international trademark clearance.
- The Marchesi and Tang evidence is a working paper.
- The legacy SUSB calculations were audited for universe meaning, not rebuilt from raw source files in this pass.
- No operator interviews or private firm financials were conducted for this draft.
- The boundary-ledger method, Dependency-Adjusted Independence assessment, and viability gate are conceptual and uncalibrated.